Rationale

The project tracks documented state-government actions involving AI, rather than the number of articles mentioning AI. An eligible record needs a named state role, an explicit AI component and evidence of an action within 1 January 2024 to 10 November 2026. Generic digitisation, private investment without a documented state role and Union-only programmes are excluded. Categories describe the action's subject; stages distinguish intention, approval, agreement, launch, pilot and operation.

Twelve states provide a purposive spread across southern, western, northern, central, eastern and northeastern India. Karnataka, Tamil Nadu, Telangana, Andhra Pradesh and Kerala offer different technology, public-service and education cases. Maharashtra and Gujarat add sector policies and administrative institutions; Uttar Pradesh, Rajasthan and Madhya Pradesh add mission, procurement and skills programmes; Odisha and Assam add gazette, procurement and compute evidence. This is not a list of the best-performing states. Availability of inspectable evidence influenced selection, and English-web coverage remains a limitation.

Official orders, budgets, gazettes and procurement documents take priority for formal decisions. Government and partner releases establish their own announcements. News supplies discovery, chronology and corroboration. Retrieved pages were reviewed against the precise claim; short excerpts, source summaries, retrieval dates and unresolved conflicts remain in a source ledger. A primary source attached to an event may support only part of it. A government-hosted newspaper clipping remains a news source.

Collection and publication are separate. Free Google News RSS and GDELT adapters collect pending candidates. Optional search adapters provide redundancy. Local HTML and PDF extraction avoids mandatory API spending; hosted extractors remain optional. Bounded requests, robots checks, retries and explicit failure logs make outages inspectable. The reviewed research files are the input to a deterministic, standard-library Python build. The snapshot is reproducible offline. The research is not presented as unattended machine extraction.

Article identity and event identity are different. Canonical URLs remove tracking noise; reviewed action identifiers join corroborating sources. Distinct counterparties and separately dated milestones remain separate. Title similarity only produces a review queue. State aliases normalize to canonical names. Event dates remain separate from publication and retrieval dates; month-only and unknown dates retain their precision. Conflicting dates and amounts are documented rather than averaged or added.

The dashboard prioritises state, category, stage, date and evidence filters. Inline source panels explain each date decision and limitation. Downloaded CSV and JSON come from the same build. A map or leaderboard would imply geographic completeness or performance comparisons the sample cannot support. GitHub Pages serves the static snapshot; scheduled Actions collect candidates without silently changing reviewed findings. The one-page rationale is supported by the methodology, research findings, source ledger, exclusion log, provider assessment and runnable tests.